

# ⌚ ASTRAL AI API — Final Architecture Summary

*Complete system overview: principles, architecture, components, and workflows*

## 📋 Table of Contents

1. [Introduction](#)
2. [High-Level Architecture](#)
3. [Core System Components](#)
4. [Crash Doctor — Crash Diagnostics System](#)
5. [Security and Limits](#)
6. [Visual Flow Diagrams](#)
7. [Technology Stack](#)
8. [Database](#)
9. [Development History](#)

## Introduction


ASTRAL AI API is an intelligent system for automatic Minecraft modpack assembly using advanced AI and RAG (Retrieval-Augmented Generation) technologies. The system analyzes user requests, finds suitable mods through semantic search, and creates ready-to-use modpack configurations.

## Key Features

- ⚡ **Conditional Architecture** — system automatically selects optimal processing path
- 🔎 **Hybrid Search** — combination of vector and keyword search
- ⚡ **Capability-Based Search** — semantic mod classification
- 🚑 **Crash Doctor** — automatic crash diagnosis and fixes
- 🔒 **Multi-Layer Protection** — JWT, rate limiting, subscription tiers
- 📊 **Full Transparency** — tracking of all stages and costs

## High-Level Architecture

# System Overview



## DEEPSEEK AI

- Query Planning
- Architecture Planning
- Mod Selection
- Crash Analysis

# Core System Components

## 1. Query Planner

**Purpose:** Analyzes user request and determines optimal processing path.

**What it does:**

- Determines request type: `simple_add`, `performance`, `themed_pack`
- Creates search plan with priorities
- Decides if Architecture Matcher is needed
- Optimizes queries for search

**Input:** User request text

**Output:** JSON search plan with metadata

```
# Example Query Planner output
{
  "request_type": "themed_pack",
  "use_architecture_matcher": true,
  "search_plan": {
    "capabilities_focus": ["combat.weapons", "building.medieval"],
    "baseline_mods": [...],
    "search_queries": [...]
  }
}
```

## 2. Hybrid Search

**Purpose:** Finds relevant mods using a combination of search methods.

**Search Methods:**

### A. Vector Search (Semantic)

- Uses embeddings (sentence-transformers)

- Understands query semantics
- Finds conceptually similar mods

## B. Keyword Search (BM25)

- Exact term matching
- Fast and efficient
- Ideal for specific mod names

## C. Reciprocal Rank Fusion

- Combines results from both methods
- Balances relevance
- Reduces bias from single method

**Input:** Search plan from Query Planner

**Output:** Ranked candidate list (80-150 mods)

---

## 3. Architecture Matcher

**Purpose:** Finds similar reference modpacks for pattern learning.

**What it does:**

- Searches database of 1000+ analyzed modpacks
- Uses vector search by theme
- Extracts capability patterns
- Identifies baseline mods

**When used:** Only for themed modpacks (50+ mods)

**Output:** Top-10 reference modpacks with their architectures

---

## 4. Architecture Planner

**Purpose:** Plans category structure for themed modpack.

**What it does:**

- Analyzes reference modpacks
- Creates 5-12 thematic categories
- Assigns required capabilities to each category
- Estimates mod count per category

**Example output:**

```
{
  "categories": [
    {
      "name": "Knight's Armory",
      "required_capabilities": ["combat.weapons", "combat.armor"],
      "target_mods": 12
    },
    {
      "name": "Castle Building",
      "required_capabilities": ["building.medieval"],
      "target_mods": 15
    }
  ]
}
```

## 5. Final Selector

**Purpose:** Intelligent selection of final mods from candidates.

**Process:**

1. **Local Pre-filtering** (100 → 50 mods)

- Filter by capabilities
- Limit by categories (max 6 mods per category)
- Remove duplicates

2. **AI Selection** (DeepSeek)

- Analyzes each mod
- Considers architecture (if available)
- Ensures category balance
- Generates explanations

**Input:** 50-100 candidates

**Output:** 20-100 selected mods with categories

---

## 6. Dependency Resolver

**Purpose:** Automatically adds all required dependencies.

**What it does:**

- Gets dependency information from database
- Recursively resolves all **required** dependencies
- Checks loader compatibility

- Handles version requirements
- Checks incompatibilities (bidirectional)

**Important:** No limit on dependency count — all required dependencies are added automatically.

**Optimization:** Batch database queries (1 query instead of 50)

---

## 7. Architecture Refiner

**Purpose:** Adjusts categories based on actually selected mods.

**What it does:**

- Analyzes real mods after selection
- Splits overloaded categories (15+ mods)
- Merges small categories (1-3 mods)
- Separates libraries from gameplay mods
- Uses thematic names

**Example:**

```
BEFORE:  
- Combat (23 mods) ← overloaded
```

```
AFTER:  
- Knight's Armory (10 mods)  
- Combat Mechanics (8 mods)  
- Battle Skills (5 mods)
```

---

## 8. Smart Categorizer

**Purpose:** Automatic categorization for simple requests.

**When used:** For Flow B (simple requests without architecture)

**What it does:**

- Analyzes description and tags of each mod
- Groups into standard categories
- Simpler than Architecture Refiner (no splitting/merging)

**Categories:** Performance, Graphics, Utility, World, Gameplay, etc.

---

## 9. Fabric Compatibility Manager

**Purpose:** Automatically handles cross-loader compatibility.

**What it does:**

- Detects when bridges between loaders are needed
- Automatically adds Connector (Fabric ↔ Forge)
- Adds Forgified Fabric API when needed
- Checks version compatibility

**Triggers:**

- Forge/NeoForge mods detected in Fabric modpack
  - `fabric_compat_mode` enabled
- 

## 10. Performance Optimizer

**Purpose:** Recommends optimization mods for specific loader.

**What it does:**


- Knows loader equivalents (Sodium → Embeddium on Forge)
  - Considers version-specific features
  - Checks optimization coverage (rendering, memory, culling)
- 

## Crash Doctor — Crash Diagnostics System

### Overview

**Crash Doctor** is a premium feature for automatic Minecraft crash diagnosis and fixes. The system analyzes crash logs and game logs, determines problem causes, and suggests specific fixes.

### Crash Doctor Architecture



# Crash Doctor Components

## 1. Log Sanitizer ( `log_sanitizer.py` )

**Purpose:** Cleans and normalizes logs before analysis.

**What it does:**

- Removes duplicate lines
- Truncates huge data blocks
- Removes PII (file paths, usernames)
- Extracts metadata (MC version, loader, mods)

**Output:** Cleaned log + metadata

---

## 2. Log Validator ( `log_validator.py` )

**Purpose:** Checks crash log freshness relative to current board\_state.

**What it does:**

- Extracts mods from crash log
- Compares with mods in board\_state
- Calculates match percentage
- Warns if log is outdated (<30% match)

**Duplicate Protection:**

- Uses MD5 hash of logs
  - 1-hour cache per user
  - Prevents re-analysis of identical logs
- 

## 3. Crash Analyzer ( `crash_analyzer.py` )

**Purpose:** Analyzes crash log using DeepSeek AI.

**What it does:**

- Reads entire crash log carefully
- Identifies ALL problem types
- Classifies errors:
  - `mod_conflict` — mod conflict
  - `missing_dependency` — missing dependency
  - `outdated_mod` — outdated mod

- o `mixin_error` — mixin transformation error
- o `class_not_found` — missing class
- o `fabric_mod_on_neoforge` — Fabric mod on NeoForge
- o `memory` — memory issues

**Output:** Structured analysis with problematic mods and causes

---

## 4. Fix Planner (`fix_planner.py`)

**Purpose:** Plans specific actions for fixes.

**Available actions:**

- `remove_mod` — remove problematic mod
- `add_mod` — add missing dependency
- `disable_mod` — disable mod
- `update_mod` — update outdated mod
- `clear_connector_cache` — clear Connector cache

**Validation:**

- Checks mod existence via Modrinth API
- Checks update availability
- Validates dependency versions

**Output:** Fix plan with priorities and justifications

---

## 5. Board Patcher (`board_patcher.py`)

**Purpose:** Creates fixed version of board\_state.

**What it does:**

- Applies operations from fix plan
- Removes/adds/updates mods
- Preserves board\_state structure
- Validates result

**Important:** Does not apply changes automatically — only creates draft

---

## 6. Session Recorder (`crash_doctor_recorder.py`)

**Purpose:** Saves each analysis session to database.

### **What is saved:**

- Unique session ID
- Processed crash log
- Analysis date
- AI solutions
- Tokens used
- Result (suggestions, patched\_board\_state)

**Goal:** Building knowledge base of solutions to improve system

---

## SSE Streaming

Crash Doctor uses Server-Sent Events for progress streaming:

### **Events:**


- `progress` — processing stage (validation, sanitization, analysis, planning, patching)
- `complete` — completion with results
- `error` — error with description

### **Advantages:**

- User sees progress in real-time
  - Long operations don't appear frozen
  - Can cancel operation if needed
- 

## Security and Limits

## Multi-Layer Protection



## 1. JWT Authentication

**Process:**

1. Client sends token in `Authorization: Bearer <token>` header
2. Server validates token via Supabase Auth API
3. Extracts `user_id` from token
4. Checks user existence in database

**Important:** Token is validated BEFORE all operations

---

## 2. Rate Limiting

**Limits by tier:**

Tier	Daily Requests	Monthly Requests	Max Mods/Request	AI Tokens/Month
<b>Free</b>	0	0	0	0
<b>Test</b>	50	1,000	50	100,000
<b>Premium</b>	200	5,000	100	500,000
<b>Pro</b>	Unlimited	Unlimited	200	Unlimited

**How it works:**

- Check BEFORE request execution

- Counters in database (`users` table)

- Automatic reset:

- Daily: every day at 00:00 UTC
- Monthly: 1st of each month

#### Errors:

- `429 Too Many Requests` — limit exceeded
  - `403 Forbidden` — Free tier blocked
- 

## 3. Subscription Tiers

#### Subscription check:


- `subscription_tier` taken ONLY from database (not from client!)
- Free users completely blocked
- Check happens via `@require_subscription` decorator


#### Protected endpoints:

- `/api/ai/build-board` — modpack assembly
  - `/api/ai/auto-sort` — automatic sorting
  - `/api/ai/crash-doctor/analyze` — crash analysis
  - `/api/feedback/ai-sort` — sorting feedback
- 

## Visual Flow Diagrams

### Flow A: Architecture-First (Themed Modpacks)





**Execution time:** ~10-15 seconds

**Cost:** ~\$0.013-\$0.026

---

## Flow B: Classic (Simple Requests)

**USER REQUEST:** "Add sodium, lithium and iris"

**Query Planner**  
• request\_type: simple  
• use\_arch\_matcher: **false**

**Hybrid Search**  
• Keyword + Semantic  
• RRF combination

**Final Selector**  
• Simple AI selection  
• **Without** architecture

**Dependency Resolver**  
• Adds dependencies


**Smart Categorizer**  
• Standard categories

**FINAL BOARD STATE**  
• **5-50** mods  
• Standard categories

**Execution time:** ~3-5 seconds

**Cost:** ~\$0.006-\$0.013

## Crash Doctor Flow



- patched\_board\_state
- confidence

**Execution time:** ~15-30 seconds

**Cost:** ~\$0.02-\$0.05 (depends on log size)

---

## Technology Stack

---

### Backend

- **Flask** — REST API framework
- **Python 3.11+** — main language
- **DeepSeek R1** — LLM for reasoning and planning
- **sentence-transformers** — local embeddings (all-MiniLM-L6-v2)
- **requests** — HTTP client for external APIs

### Database

- **Supabase** — PostgreSQL + pgvector
- **pgvector** — vector search
- **PostgreSQL** — relational database

### External APIs

- **Modrinth API** — mod, version, dependency information
- **Supabase Auth API** — JWT validation
- **DeepSeek API** — AI requests

### Infrastructure

- **Cloudflare Tunnel** — tunneling for production
  - **SSE (Server-Sent Events)** — progress streaming
- 

### Database

---

#### Core Tables

## mods

Stores information about all mods.

### Key fields:

- `id` (UUID) — unique ID
- `source_id` (TEXT) — Modrinth project ID
- `embedding` (VECTOR(384)) — vector representation for search
- `capabilities` (TEXT[]) — AI-generated capabilities
- `tags` (TEXT[]) — custom tags
- `loaders` (TEXT[]) — supported loaders
- `game_versions` (TEXT[]) — supported MC versions
- `dependencies` (JSONB) — dependencies
- `incompatibilities` (JSONB) — incompatibilities

### Indexes:

- `idx_embedding` — for vector search (ivfflat)
- 

## modpacks

Reference modpacks for pattern learning.

### Key fields:

- `id` (UUID)
- `title`, `slug`, `description`
- `architecture` (JSONB) — capability and provider structure
- `embedding` (VECTOR(384)) — for semantic search
- `mc_version`, `mod_loader`

**Usage:** Architecture Matcher uses this table to find similar modpacks

---

## users

User information and subscriptions.

### Key fields:

- `id` (UUID) — matches user\_id from JWT
- `subscription_tier` (TEXT) — 'free', 'test', 'premium', 'pro'
- `daily_requests_used` (INTEGER) — requests used today
- `monthly_requests_used` (INTEGER) — requests used this month

- `ai_tokens_used` (INTEGER) — AI tokens used
  - `last_request_date` (DATE) — last request date
  - `custom_limits` (JSONB) — custom limits for VIP
- 

### `modpack_builds`

Saves all AI-generated modpacks for learning.

**Key fields:**

- `id` (TEXT) — unique ID (format: 0000001)
- `user_id` (UUID) — user
- `prompt` (TEXT) — original request
- `architecture` (JSONB) — generated architecture
- `feedback` (JSONB) — user feedback
- `learning_status` (JSONB) — feedback processing status

**Usage:** System learning based on successful builds

---

### `ai_sort_sessions`

Tracks automatic sorting sessions.

**Key fields:**

- `id` (TEXT) — unique ID
- `user_id` (UUID)
- `input_mods` (TEXT[]) — mods before sorting
- `categories` (JSONB) — generated categories
- `rating` (INTEGER) — user rating (1-5 stars)
- `created_at` (TIMESTAMPTZ)

**Usage:** Improving categorization quality

---

### `crash_doctor_sessions`

Crash Doctor solution knowledge base.

**Key fields:**

- `id` (TEXT) — unique session ID
- `user_id` (UUID)

- `crash_log` (TEXT) — processed crash log
- `board_state` (JSONB) — board state at analysis time
- `suggestions` (JSONB) — suggested fixes
- `patched_board_state` (JSONB) — fixed state
- `confidence` (FLOAT) — solution confidence
- `token_usage` (JSONB) — tokens used
- `created_at` (TIMESTAMPTZ)

**Usage:** Building knowledge base to improve diagnostics

---

## Development History

---

### Version 1.0 (Legacy)

- Basic AI mod search
- Simple categorization
- Minimal validation

### Version 2.0

- Added Hybrid Search
- Improved categorization
- Added dependency support

### Version 3.0 (Current) — Conditional Architecture

- **Conditional Architecture** — automatic flow selection
- **Architecture Matcher** — reference modpack search
- **Architecture Planner** — category planning
- **Architecture Refiner** — category refinement
- **Capability-based matching** — capability search
- **Baseline mods** — automatic inclusion of base mods

### Crash Doctor (Added later)

- Full crash analysis pipeline
- SSE streaming for progress
- Log validation
- Duplicate caching
- Solution knowledge base

## Security (Continuously improved)

- JWT authentication
  - Rate limiting by tier
  - Subscription tier protection
  - Request logging
- 

## Key Design Principles

### 1. Conditional Architecture

**Problem:** One pipeline doesn't work well for all request types.

**Solution:** Automatic selection of optimal path based on request analysis.

**Advantages:**

- ⚡ Fast for simple requests
  - 🎯 High quality for complex requests
  - 💰 Cost efficiency
- 

### 2. Capabilities over Tags

**Problem:** Tags are too broad and don't allow precise matching.

**Solution:** AI-generated capabilities with hierarchical structure.

**Advantages:**

- More granular classification (50+ capabilities vs 20-30 tags)
  - Hierarchical structure (combat.weapons.melee)
  - Precise matching for categorization
- 

### 3. Reference Modpack Learning

**Idea:** Popular modpacks represent proven combinations.

**Approach:**

1. Parsing 1000+ successful modpacks from Modrinth
2. Extracting capability patterns
3. Using as templates for new builds

### **Advantages:**

- Learning from community experience
  - Avoiding common incompatibilities
  - Ensuring base quality
- 

## **4. Pipeline Transparency**

**Principle:** Full traceability of all operations.

### **What is tracked:**

- AI token usage
- API request costs
- Search and ranking scores
- Mod selection justifications
- Execution time for each stage

### **Advantages:**

- Debugging and optimization
  - Cost monitoring
  - Quality analysis
- 

## **Metrics and Performance**

### **Execution Time**

- **Simple requests** (Flow B): 3-5 seconds
- **Themed modpacks** (Flow A): 10-15 seconds
- **Crash Doctor**: 15-30 seconds

### **Cost (DeepSeek API)**

- **Simple requests**: ~\$0.006-\$0.013
- **Themed modpacks**: ~\$0.013-\$0.026
- **Crash Doctor**: ~\$0.02-\$0.05

### **Optimizations**

- **Batch DB queries** — 80% reduction in dependency resolution time
- **Local prefILTERING** — reduces AI candidates from 100 to 50

- **Capability-based matching** — more accurate search
  - **Log caching** — prevents duplicate analysis
- 

## Future Improvements

---

### Planned Features

1. **Semantic Clusters**
    - Pre-built "skeletons" for common modpack types
    - Fast assembly for known themes
  2. **Co-occurrence Graph**
    - "Mods that work well together"
    - Learning from successful modpack combinations
  3. **User Preference Learning**
    - Tracking user build history
    - Personalized recommendations
  4. **Multi-stage Architecture Planning**
    - Iterative refinement with user feedback
    - Interactive category configuration
- 

## Conclusion

---

ASTRAL AI API is a comprehensive system for intelligent Minecraft modpack assembly, combining:

-  **Advanced AI technologies** (DeepSeek R1)
-  **Hybrid search** (vector + keyword)
-  **Semantic understanding** (capabilities)
-  **Automatic diagnostics** (Crash Doctor)
-  **Multi-layer protection** (JWT + Rate Limiting)
-  **Full transparency** (tracking all operations)

The system continuously evolves, learning from user feedback and improving recommendation quality.

---

**Document Version:** 1.0

**Last Updated:** 2025-11-12

**Author:** ASTRAL Team

---

*Built with ❤ for the Minecraft modding community*